



고려대학교
KOREA UNIVERSITY

KU-The Future

Cloud computing – Distributed Systems

Heonchang Yu

Distributed and Cloud Computing Lab.

Clocks, events and process states

– Distributed system

- A collection P of N processes $p_i, i = 1, 2, \dots, N$
- Each process p_i has a state s_i consisting of its variables (which it transforms as it executes)
- Processes communicate only by messages (via a network)
- Actions of processes
 - ✓ a message Send or Receive operation
 - ✓ an operation that transforms process' state
- Event - the occurrence of a single action that a process carries out as it executes <Ex.> Send, Receive, state-transforming action
- The sequence of events within a single process p_i can be placed in a single, total ordering, denoted by the relation \rightarrow_i between the events. i.e. $e \rightarrow_i e'$ if and only if the event e occurs before e' at p_i
- A history of process p_i is a series of events ordered by \rightarrow_i
 $history(p_i) = h_i = \langle e_i^0, e_i^1, e_i^2, \dots \rangle$

Global states

- Distributed garbage collection
 - To be garbage if there are no longer any references to it anywhere in the distributed system
 - To check that an object is garbage, we must verify that there are no references to it anywhere in the system
 - Need to check a reference to it in a message that is in transit between the processes
- Distributed deadlock detection
 - There is a cycle in the graph of this 'waits-for' relationship.
- Distributed termination detection
 - To test whether each process has halted

Global states

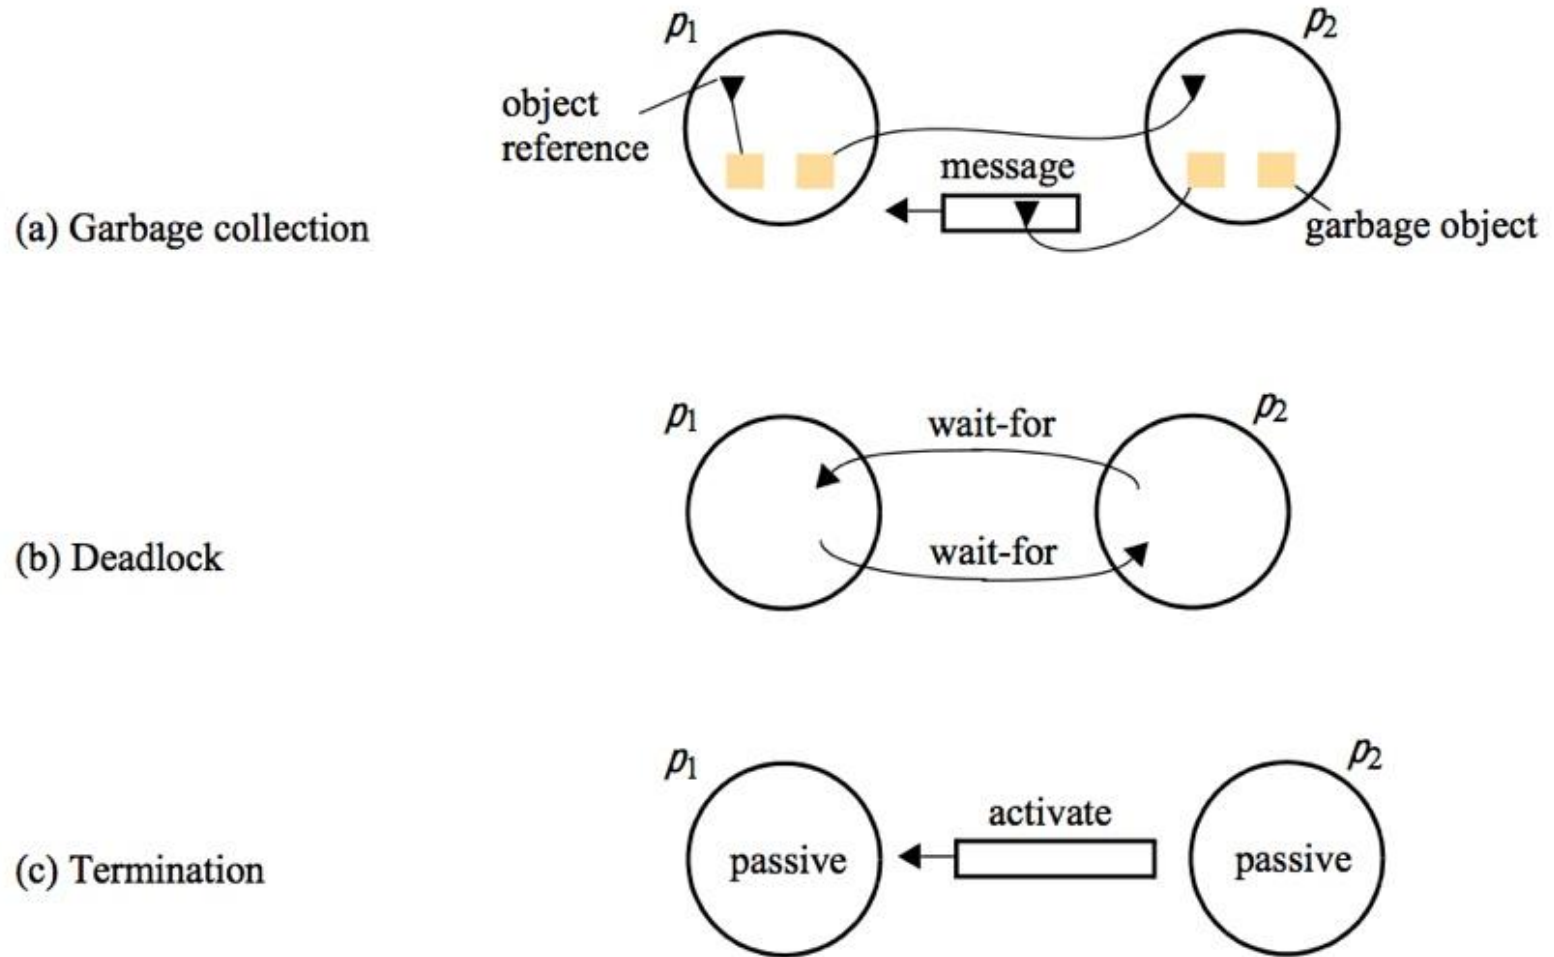


Figure 14.8 Detecting global properties

Global states

- Global states and consistent cuts
 - Problem: absence of global time
 - Event : an internal action of the process or the sending or receipt of a message over the communication channels
 - Cut - A subset of its global history
 - ✓ Take any set of states of the individual processes to form a global state $S=(s_1, s_2, \dots, s_N)$
 - ✓ At p_2 it includes the receipt of the message m_1 , but at p_1 it does not include the sending of that message – showing an 'effect' without a 'cause'
 - Consistent cut - if, for each event it contains, it also contains all the events that happened-before that event:
for all events $e \in C, f \rightarrow e \Rightarrow f \in C$
 - Consistent global state is one that corresponds to a consistent cut.

Global states




Figure 14.9 Cuts

Global states

- Global state predicates, stability, safety and liveness
 - Stable: once the system enters a state in which the predicates associated is *True*, it remains *True* in all future states reachable from that state.
 - Safety: the assertion that α (being deadlocked) evaluates to *False* for all states S reachable from S_0
 - Liveness: the property that, for any linearization L starting in the state S_0 , β evaluates to *True* for some state S_L reachable from S_0

Global states

- The 'snapshot' algorithm of Chandy and Lamport
 - Determining global states of distributed systems
 - Assumptions
 - ✓ Neither channels nor processes fail; communication is reliable so that every message sent is eventually received intact, exactly once;
 - ✓ Channels are unidirectional and provide FIFO-ordered message delivery;
 - ✓ The graph of processes and channels is strongly connected;
 - ✓ Any process may initiate a global snapshot at any time;
 - ✓ The processes may continue their execution and send and receive normal messages while the snapshot takes place.
 - System model
 - ✓ Incoming and Outgoing channels
 - ✓ Special marker messages
 - ❖ A prompt for a receiver to save its own state, if it has not already done so
 - ❖ A means of determining which messages to include in the channel state

Global states

– The 'snapshot' algorithm of Chandy and Lamport

- Two rules of algorithm


- ✓ Marker receiving rule

- ❖ Obligates a process that has not recorded its state to do so
- ❖ When a process that has already saved its state receives a marker, it records the state of that channel as the set of messages it received on it since it saved on it.

- ✓ Marker sending rule

- ❖ Obligates processes to send a marker after they have recorded their state, but before they send any other messages

- ✓ Example



Global states

Marker receiving rule for process p_i

On p_i 's receipt of a *marker* message over channel c :

if (p_i has not yet recorded its state) it

records its process state now;

records the state of c as the empty set;

turns on recording of messages arriving over other incoming channels;

else

p_i records the state of c as the set of messages it has received over c
since it saved its state.

end if

Marker sending rule for process p_i

After p_i has recorded its state, for each outgoing channel c :

p_i sends one marker message over c

(before it sends any other message over c).

Figure 14.10 Chandy and Lamport's 'snapshot' algorithm

Global states

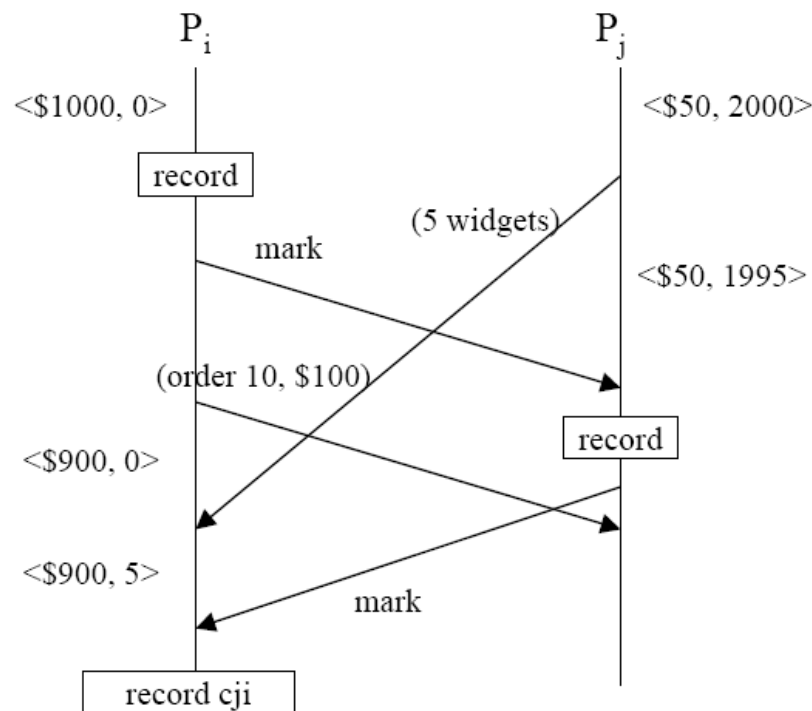
- The 'snapshot' algorithm of Chandy and Lamport
 - Example: a system of two processes, p_1 and p_2 connected by two unidirectional channels, c_1 and c_2
 - ✓ Two processes trade in 'widgets'.
 - ✓ Process p_1 sends orders for widgets over c_2 to p_2 enclosing payment at the rate of \$10 per widget.
 - ✓ Process p_2 sends widgets along channel c_1 to p_1 .
 - ✓ Process p_2 has already received an order for five widgets
 - ❖ S_0 : Process p_1 records its state when p_1 's state is $\langle \$1000, 0 \rangle$.
 - ❖ S_1 : Following the marker sending rule, process p_1 emits a marker message over its outgoing channel c_2 before it sends the next application-level message: (order 10, \$100) over channel c_2 .
 - ❖ S_2 : Before p_2 receives the marker, it emits an application message (five widgets) over c_1 in response to p_1 's previous order. Process p_1 receives p_2 's message (five widgets), and p_2 receives the marker. Following the marker receiving rule, p_2 records its state as $\langle \$50, 1995 \rangle$ and that of channel c_2 as the empty sequence. Following the marker sending rule, it sends a marker message over c_1 .

Global states

– The 'snapshot' algorithm of Chandy and Lamport

❖ S_3 : When process p_1 receives p_2 's marker message, it records the state of channel c_1 as the single message (five widgets) that it received after it first recorded its state.

✓ $p_1 : \langle \$1000, 0 \rangle$; $p_2 : \langle \$50, 1995 \rangle$; $c_1 : \langle \text{(five widgets)} \rangle$; $c_2 : \langle \rangle$



Logical time and logical clocks

- As Lamport pointed out, since we cannot synchronize clocks perfectly across a distributed system, we cannot use physical time to find out the order of any arbitrary pair of events.
- For ordering
 - ✓ If two events occurred at the same process p_i ($i=1,\dots,N$), then they occurred in the order in which p_i observes them.
 - ✓ Whenever a message is sent between processes, the event of sending the message occurred before the event of receiving the message.
- Happened-before relation (\rightarrow)
 - ✓ HB1: If \exists process $p_i: e \rightarrow_i e'$, then $e \rightarrow e'$.
 - ✓ HB2: For any message m , $send(m) \rightarrow receive(m)$
 - where $send(m)$ is the event of sending the message, and $receive(m)$ is the event of receiving it.
 - ✓ HB3: If e , e' and e'' are events such that $e \rightarrow e'$ and $e' \rightarrow e''$, then $e \rightarrow e''$.

Logical time and logical clocks

- In figure 14.5, not all events are related by the relation \rightarrow .
 - $a \not\rightarrow e$ and $e \not\rightarrow a$
 - ✓ Since they occur at different processes, and there is no chain of messages intervening between them.
 - ✓ Concurrent : $a \parallel e$




Figure 14.5 Events occurring at three processes

Logical time and logical clocks

- Logical clocks: a monotonically increasing software counter
 - LC1: L_i is incremented before each event is issued at process p_i :
$$L_i := L_i + 1$$
 - LC2: (a) When a process p_i sends a message m , it piggybacks on m the value $t = L_i$
(b) On receiving (m, t) , a process p_j computes $L_j := \max(L_j, t)$ and then applies LC1 before timestamping the event $receive(m)$.




Figure 14.6 Lamport timestamps for the events shown in Figure 14.5

Logical time and logical clocks

– Logical clocks

- For two events e and e' , $e \rightarrow e' \Rightarrow L(e) < L(e')$.
- The converse is not true: $L(e) < L(e') \nRightarrow e \rightarrow e'$.
✓ <ex> $L(b) > L(e)$ but $b \parallel e$.

– Totally ordered logical clocks

- If e is an event occurring at p_i with local timestamp T_i , and e' is an event occurring at p_j with local timestamp T_j , we define the global logical timestamps for these events to be (T_i, i) and (T_j, j) .
✓ $(T_i, i) < (T_j, j) : T_i < T_j$, or $T_i = T_j$ and $i < j$

– Vector clocks

- To overcome the shortcoming of Lamport's clocks: the fact that from $L(e) < L(e')$ conclude that $e \rightarrow e'$
- A vector clock for a system of N processes is an array of N integers.
- Processes piggyback vector timestamps on the messages they send to one another.

Logical time and logical clocks

– Rules of vector clocks

- VC1: Initially, $V_i[j] = 0$, for $i, j=1, 2, \dots, N$
- VC2: Just before p_i timestamps an event, it sets $V_i[i] := V_i[i] + 1$.
- VC3: p_i includes the value $t = V_i$ in every message it sends.
- VC4: When p_i receives a timestamp t in a message, it sets $V_i[j] := \max(V_i[j], t[j])$, for $j=1, 2, \dots, N$. Taking the component-wise maximum of two vector timestamps in this way is known as a merge operation.




Figure 14.7 Vector timestamps for the events shown in Figure 14.5