



고려대학교
KOREA UNIVERSITY

KU-The Future

Cloud computing – Distributed Systems

Heonchang Yu

Distributed and Cloud Computing Lab.

Distributed deadlocks

- **Deadlocks**

- ✓ Can arise within a single server when locking is used for concurrency control
- ✓ Servers must either prevent or detect and resolve **deadlocks**.
- ✓ Using **timeouts** to resolve possible deadlocks is a clumsy approach.
- ✓ Most **deadlock detection schemes** operate by finding cycles in the transaction wait-for graph.
 - ❖ Nodes: transactions and objects
 - ❖ Edges: an object held by a transaction or a transaction waiting for an object
 - ❖ **Distributed deadlock**: can be a cycle in the global wait-for graph that is not in any single local one
- ✓ In figure 1(table) and figure 2(graph)
 - ❖ The objects A and B managed by servers X and Y and objects C and D managed by server Z
 - ❖ A deadlock cycle in figure 2(a)
 - ❖ Wait-for graph in figure 2(b)

Distributed deadlocks

<i>U</i>	<i>V</i>	<i>W</i>
<i>d.deposit(10)</i> lock <i>D</i>		
<i>a.deposit(20)</i> lock <i>A</i> at <i>X</i>	<i>b.deposit(10)</i> lock <i>B</i> at <i>Y</i>	
<i>b.withdraw(30)</i> wait at <i>Y</i>		<i>c.deposit(30)</i> lock <i>C</i> at <i>Z</i>
	<i>c.withdraw(20)</i> wait at <i>Z</i>	<i>a.withdraw(20)</i> wait at <i>X</i>

Figure 1 Interleaving of transaction *U*, *V* and *W*

Distributed deadlocks

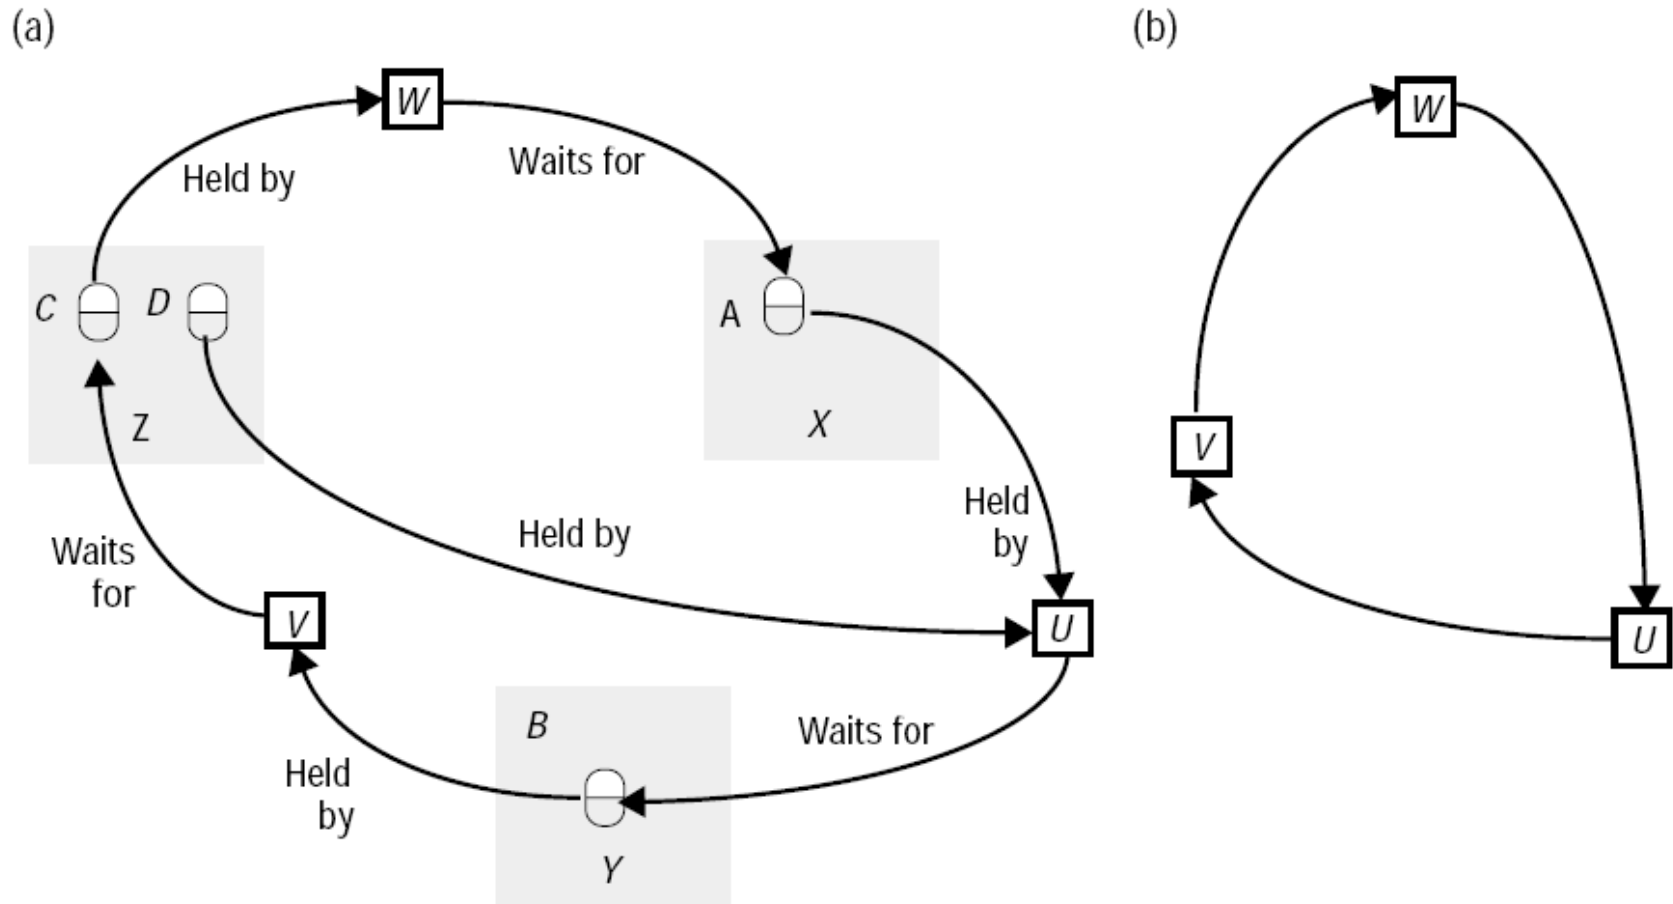


Figure 2 Distributed deadlock

Distributed deadlocks

- **Deadlocks**

- ✓ Local wait-for graphs can be built by the lock manager at each server.

- ❖ Server Y : $U \rightarrow V$ (added when U requests $b.withdraw(30)$)
 - ❖ Server Z : $V \rightarrow W$ (added when V requests $c.withdraw(20)$)
 - ❖ Server X : $W \rightarrow U$ (added when W requests $a.withdraw(20)$)

- ✓ **Centralized deadlock detection**


- ❖ **One server** takes on the role of global deadlock detector.
 - ❖ Each server sends the latest copy of its **local wait-for graph** to the global deadlock detector.
 - ❖ When the detector finds a cycle, it **makes a decision** on how to resolve the deadlock and **informs the servers** as to the transaction to be aborted to resolve the deadlock.
 - ❖ Centralized deadlock detection is **not a good idea**, because it depends on a single server to carry it out.
 - ❖ **Problems:** poor availability, lack of fault tolerance, no ability to scale and the cost of the frequent transmission of local wait-for graph

Distributed deadlocks


- Phantom deadlocks

- ✓ A deadlock that is 'detected' but is **not really a deadlock**
- ✓ As the procedure for deadlock detection will **take some time**, there is a chance that one of the transactions that holds a lock will **meanwhile have released** it, in which case the deadlock will no longer exist.
- ✓ In figure 3, it would detect a cycle $T \rightarrow U \rightarrow V \rightarrow T$, although the edge $T \rightarrow U$ **no longer exists**.
 - ❖ If there is a cycle $T \rightarrow U \rightarrow V \rightarrow T$ and U aborts after the information concerning U has been collected, then the cycle has been already and there is no deadlock.

local wait-for graph



local wait-for graph



global deadlock detector




Figure 3 Local and global wait-for graphs

Distributed deadlocks

- Edge chasing(=path pushing)

- ✓ The global wait-for graph is not constructed, but each of the servers involved has **knowledge** about some of its edges.
- ✓ The servers attempt to find cycles by forwarding **messages called probes**.
- ✓ A probe message consists of **transaction wait-for relationships** representing a path in the global wait-for graph.
- ✓ The situation at server X in **figure 2**
 - ❖ The server X has just added **the edge $W \rightarrow U$** to its local wait-for graph and at this time, transaction **U is waiting to access object B** , which transaction V holds at server Y .
 - ❖ This edge could **possibly be part of a cycle** such as $V \rightarrow T_1 \rightarrow T_2 \rightarrow \dots \rightarrow W \rightarrow U \rightarrow V$ involving transactions using objects at other servers.
 - ❖ This indicates that there is **a potential distributed deadlock cycle**, which could be found by sending out a probe to server Y .

Distributed deadlocks

- Edge chasing algorithm

- ❖ **Initiation:** When a server notes that a **transaction T starts waiting for another transaction U** , where U is waiting to access an object at another server, it initiates detection by sending a probe containing the **edge $\langle T \rightarrow U \rangle$** to the server of the object at which transaction U is blocked. If U is sharing a lock, probes are sent to all the holders of the lock.
- ❖ **Detection:** Detection consists of **receiving** probes and **deciding** whether deadlock has occurred and whether to forward the probes.
 - When a server of an object receives a probe $\langle T \rightarrow U \rangle$, it checks to see whether U is also waiting.
 - If it is, the transaction it waits for (for example, V) is added to the probe (making it $\langle T \rightarrow U \rightarrow V \rangle$), and if the new transaction (V) is waiting for another object elsewhere, the probe is forwarded.
 - Before forwarding a probe, the server checks to see whether the transaction it has just added has caused the probe to contain a cycle (ex. $\langle T \rightarrow U \rightarrow V \rightarrow T \rangle$). If this is the case, it has found a cycle in the graph and deadlock has been detected.
- ❖ **Resolution:** When a cycle is detected, a transaction in the cycle is **aborted** to break the deadlock.

Distributed deadlocks

- Edge chasing

- ✓ The example of figure 4

- ❖ **Server X** initiates detection by sending **probe** $\langle W \rightarrow U \rangle$ to the server of B (Server Y)

- ❖ **Server Y** receives probe $\langle W \rightarrow U \rangle$, notes that B is held by V and appends V to the **probe to produce** $\langle W \rightarrow U \rightarrow V \rangle$. It notes that V is waiting for C at server Z . This probe is forwarded to server Z .

- ❖ **Server Z** receives probe $\langle W \rightarrow U \rightarrow V \rangle$ and notes C is held by W and appends W to the **probe to produce** $\langle W \rightarrow U \rightarrow V \rightarrow W \rangle$.

- ✓ A probe that detects a cycle involving N transactions will be forwarded **by $(N-1)$ transaction coordinators** via **$(N-1)$ servers** of objects, requiring $2(N-1)$ messages.

Distributed deadlocks




Figure 4 Probes transmitted to detect deadlock

Distributed deadlocks

- **Transaction priorities**

- ✓ The effect of several transactions in a cycle initiating deadlock detection is that detection may **happen at several different servers** in the cycle with the result that more than one transaction in the cycle is aborted.
- ✓ In figure 5
 - ❖ (a) : U is waiting for W and V is waiting for T . T requests the object held by U and W requests the object held by V . **Two separate probes $\langle T \rightarrow U \rangle$ and $\langle W \rightarrow V \rangle$** are initiated by the servers of these objects and are circulated until deadlock is detected by each of two different servers.
 - ❖ (b), (c) : Cycles are $\langle T \rightarrow U \rightarrow W \rightarrow V \rightarrow T \rangle$ and $\langle W \rightarrow V \rightarrow T \rightarrow U \rightarrow W \rangle$.
- ✓ In order to ensure that only one transaction in a cycle is aborted, transactions are given **priorities** in such a way that all transactions are totally ordered. **(ex. Timestamps)**
 - ❖ When a deadlock cycle is found, the transaction with the **lowest priority is aborted.**
 - ❖ In the example of figure 5, **assume $T > U > V > W$.** Then the transaction W will be aborted when a cycle is detected.

Distributed deadlocks

- Transaction priorities

- ✓ Transaction priorities could also be used **to reduce** the number of situations that cause deadlock detection to be initiated.
- ✓ Transaction priorities could also be used **to reduce** the number of probes that are forwarded.
 - ❖ Probes should travel '**downhill**' – that is, from transactions with higher priorities to transactions with lower priorities.
 - ❖ Servers use the rule that they **do not forward** any probe to a holder that has higher priority than the initiator.

(a) initial situation

(b) detection initiated at object requested by T

(c) detection initiated at object requested by W




Figure 5 Two probes initiated

Distributed deadlocks

- ✓ Without priority rules, detection is initiated when W starts waiting by sending a probe $\langle W \rightarrow U \rangle$.
- ✓ Under the priority rule, this probe will not be sent, because $W < U$ and deadlock will not be detected.
- ✓ In figure 6
 - ❖ (a) : When U starts waiting for V , the coordinator of V will save the probe $\langle U \rightarrow V \rangle$
 - ❖ (b) : When V starts waiting for W , the coordinator of W will store $\langle V \rightarrow W \rangle$ and V will forward its probe queue, $\langle U \rightarrow V \rangle$, to W .
 - ❖ W 's probe queue has $\langle U \rightarrow V \rangle$ and $\langle V \rightarrow W \rangle$. When W starts waiting for A it will forward its probe queue $\langle U \rightarrow V \rightarrow W \rangle$ to the server of A , which also notes the new dependency $W \rightarrow U$ and combines it with the information in the probe received to determine that $U \rightarrow V \rightarrow W \rightarrow U$. Deadlock is detected.

(a) V stores probe when U starts waiting (b) Probe is forwarded when V starts waiting




Figure 6
Probes travel downhill